Assignment 3 — Deep Q-Network (DQN) on Atari Pong

Course: CSCN 8020 — Reinforcement Learning Programming
Submitted By : Aiswarya Thekkuveettil Thazhath - 8993970

Overview

In this assignment we implement the **Deep Q-Network (DQN)** algorithm to train an agent to play the Atari game **Pong** (`PongDeterministic-v4`). Because the state space is continuous (raw pixel frames), a tabular Q-Learning approach is infeasible. Instead, we approximate the Q-function with a **Convolutional Neural Network (CNN)** that takes the last **4 preprocessed game frames** stacked together as input.

Key Design Choices

Component	Choice
State representation	4 stacked grayscale frames (84 × 80 × 4)
Q-function approximator	CNN (3 conv layers + 2 FC layers)
Replay memory	Fixed-size deque, uniform sampling
Target network	Separate network, updated every N episodes
Exploration	ε-greedy with exponential decay

1. Environment Setup & Imports

We import all required libraries. The `assignment3_utils.py` helper provided by the professor contains the image pre-processing functions we will use throughout.

In [1]:
```
# ── Standard Library ──────────────────────────────────────────────
import os
import random
import time
from collections import deque
from typing import Tuple, List, Deque

# ── Scientific stack ──────────────────────────────────────────────
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
matplotlib.rcParams['figure.dpi'] = 120

# ── Deep Learning ─────────────────────────────────────────────────
import torch
import torch.nn as nn
import torch.optim as optim
```

```python
import torch.nn.functional as F

# ── Gym / Atari ──────────────────────────────────────────────────────────
import gymnasium as gym
import ale_py
gym.register_envs(ale_py)

# ── Professor-provided utilities ─────────────────────────────────────────
from assignment3_utils import process_frame, transform_reward

# ── Reproducibility ──────────────────────────────────────────────────────
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {DEVICE}")
```

```
Using device: cuda
```

In [2]:
```python
# Test cell - run this to confirm ALE is registered
env = gym.make("ALE/Pong-v5", render_mode=None)
print("Success!")
env.close()
```

```
Success!
```

2. Hyperparameters

All hyperparameters are centralised in a single dictionary so they can be swapped easily for the ablation experiments later.

In [3]:
```python
# ── Default hyperparameters (as specified in the assignment) ─────────────
DEFAULT_HPARAMS = {
    # --- Environment ---
    "env_name" : "ALE/Pong-v5",
    "n_actions"        : 6,
    # Image shape AFTER crop + downsample  →  84 rows × 80 cols × 1 channel
    "image_shape"      : (84, 80),
    "n_frames"         : 4,               # Stack the last 4 frames

    # --- Training ---
    "n_episodes"       : 400,             # Total training episodes
    "max_steps"        : 10_000,          # Safety cap per episode
    "batch_size"       : 8,               # Mini-batch size  ← assignment default
    "target_update_ep" : 10,              # Target-net update frequency (episodes)  ← default
    "memory_size"      : 10_000,          # Replay buffer capacity
    "min_memory"       : 500,             # Minimum memories before training starts

    # --- Q-Learning ---
    "gamma"            : 0.95,            # Discount factor γ
    "lr"               : 1e-4,            # Adam learning-rate

    # --- Exploration (ε-greedy) ---
    "epsilon_init"     : 1.0,
    "epsilon_decay"    : 0.995,           # δ
    "epsilon_min"      : 0.05,
```

```
}

print("Default hyperparameters loaded.")
for k, v in DEFAULT_HPARAMS.items():
    print(f"  {k:20s} = {v}")
```

```
Default hyperparameters loaded.
  env_name             = ALE/Pong-v5
  n_actions            = 6
  image_shape          = (84, 80)
  n_frames             = 4
  n_episodes           = 400
  max_steps            = 10000
  batch_size           = 8
  target_update_ep     = 10
  memory_size          = 10000
  min_memory           = 500
  gamma                = 0.95
  lr                   = 0.0001
  epsilon_init         = 1.0
  epsilon_decay        = 0.995
  epsilon_min          = 0.05
```

3. CNN Architecture — `DQNNetwork`

The CNN receives a tensor of shape **(batch, 4, 84, 80)** — the 4 stacked frames are treated as 4 input channels (instead of blending them together). This lets the network learn temporal motion between frames.

Architecture

Layer	Type	Kernel	Filters	Activation
Conv1	Conv2d	8×8, stride 4	32	ReLU
Conv2	Conv2d	4×4, stride 2	64	ReLU
Conv3	Conv2d	3×3, stride 1	64	ReLU
Flatten	—	—	—	—
FC1	Linear	—	512	ReLU
FC2	Linear	—	n_actions	— (linear output)

In [4]:
```
class DQNNetwork(nn.Module):
    """
    Deep Q-Network approximating Q(s, a) for all actions simultaneously.

    Input shape : (batch_size, n_frames, H, W)  →  (B, 4, 84, 80)
    Output shape: (batch_size, n_actions)        →  (B, 6)
    """

    def __init__(self, n_frames: int, n_actions: int, image_shape: Tuple[int, int]) -> None:
        super(DQNNetwork, self).__init__()
        self.n_frames  = n_frames
        self.n_actions = n_actions
        self.H, self.W = image_shape
```

```python
        # ── Convolutional backbone ──────────────────────────────────────
        self.conv_layers = nn.Sequential(
            nn.Conv2d(in_channels=n_frames, out_channels=32,
                      kernel_size=8, stride=4, padding=0),   # → (32, 20, 19)
            nn.ReLU(),
            nn.Conv2d(in_channels=32, out_channels=64,
                      kernel_size=4, stride=2, padding=0),   # → (64,  9,  8)
            nn.ReLU(),
            nn.Conv2d(in_channels=64, out_channels=64,
                      kernel_size=3, stride=1, padding=0),   # → (64,  7,  6)
            nn.ReLU(),
        )

        # Dynamically compute the flattened size after convolutions
        conv_out_size = self._get_conv_output_size()

        # ── Fully-connected head ────────────────────────────────────────
        self.fc_layers = nn.Sequential(
            nn.Linear(conv_out_size, 512),
            nn.ReLU(),
            nn.Linear(512, n_actions),   # One Q-value per action (no activation)
        )

    # ── Helper: compute conv output size without a forward pass on real data ──
    def _get_conv_output_size(self) -> int:
        dummy = torch.zeros(1, self.n_frames, self.H, self.W)
        out   = self.conv_layers(dummy)
        return int(np.prod(out.shape[1:]))   # channel × height × width

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Forward pass.

        Parameters
        ----------
        x : torch.Tensor  shape (B, n_frames, H, W),  dtype float32, range [-1, 1]

        Returns
        -------
        torch.Tensor  shape (B, n_actions) — raw Q-values
        """
        features = self.conv_layers(x)          # (B, 64, 7, 6)
        features = features.flatten(start_dim=1)  # (B, flat)
        return self.fc_layers(features)           # (B, n_actions)


# ── Sanity-check the architecture ──────────────────────────────────────
_hp   = DEFAULT_HPARAMS
_net  = DQNNetwork(_hp["n_frames"], _hp["n_actions"], _hp["image_shape"]).to(DEVICE)
_dummy = torch.zeros(1, _hp["n_frames"], *_hp["image_shape"]).to(DEVICE)
_out  = _net(_dummy)
print(f"Network output shape: {_out.shape}  (expected [1, {_hp['n_actions']}])")
print()
print(_net)
```

```
Network output shape: torch.Size([1, 6])  (expected [1, 6])

DQNNetwork(
  (conv_layers): Sequential(
    (0): Conv2d(4, 32, kernel_size=(8, 8), stride=(4, 4))
    (1): ReLU()
    (2): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2))
    (3): ReLU()
    (4): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))
    (5): ReLU()
  )
  (fc_layers): Sequential(
    (0): Linear(in_features=2688, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=6, bias=True)
  )
)
```

4. Replay Memory — `ReplayBuffer`

Experience replay breaks the temporal correlation between consecutive training samples by storing `(state, action, reward, next_state, done)` transitions and sampling uniformly at random.

In [5]:
```python
class ReplayBuffer:
    """
    Fixed-size circular replay buffer.

    Stores experience tuples: (state, action, reward, next_state, done)
    and supports uniform random sampling for mini-batch training.
    """

    def __init__(self, capacity: int) -> None:
        self.buffer: Deque = deque(maxlen=capacity)

    # ── Public interface ──────────────────────────────────────────────
    def push(self,
             state:      np.ndarray,
             action:     int,
             reward:     float,
             next_state: np.ndarray,
             done:       bool) -> None:
        """Store one transition."""
        self.buffer.append((state, action, reward, next_state, done))

    def sample(self, batch_size: int) -> Tuple:
        """
        Sample a random mini-batch.

        Returns
        -------
        Tuple of np.ndarrays:
            states, actions, rewards, next_states, dones
        """
        batch = random.sample(self.buffer, batch_size)
        states, actions, rewards, next_states, dones = zip(*batch)
        return (
            np.array(states,      dtype=np.float32),
```

```
            np.array(actions,     dtype=np.int64),
            np.array(rewards,     dtype=np.float32),
            np.array(next_states, dtype=np.float32),
            np.array(dones,       dtype=np.float32),
        )

    def __len__(self) -> int:
        return len(self.buffer)


print("ReplayBuffer class defined.")
```

```
ReplayBuffer class defined.
```

5. DQN Agent — `DQNAgent`

The agent encapsulates:

- The **online network** (trained every step)
- The **target network** (updated every `target_update_ep` episodes)
- The **replay buffer**
- The **ε-greedy policy**
- The **learn** method implementing the Bellman target update

In [6]:
```
class DQNAgent:
    """
    Deep Q-Network agent.

    Maintains two networks (online + target) and a replay buffer.
    Implements ε-greedy exploration with exponential decay.
    """

    def __init__(self, hparams: dict, device: torch.device = DEVICE) -> None:
        self.hp     = hparams
        self.device = device

        # ── Networks ──────────────────────────────────────────────────────
        self.online_net = DQNNetwork(
            hparams["n_frames"], hparams["n_actions"], hparams["image_shape"]
        ).to(device)

        self.target_net = DQNNetwork(
            hparams["n_frames"], hparams["n_actions"], hparams["image_shape"]
        ).to(device)

        # Initialise target network with the same weights as online network
        self._sync_target_network()

        # Target network is never directly trained — gradients not needed
        for p in self.target_net.parameters():
            p.requires_grad = False

        # ── Optimiser & loss ──────────────────────────────────────────────
        self.optimizer = optim.Adam(self.online_net.parameters(), lr=hparams["lr"])
        self.loss_fn   = nn.MSELoss()
```

```python
        # ── Replay buffer ──────────────────────────────────────────────
        self.memory = ReplayBuffer(hparams["memory_size"])

        # ── Exploration ────────────────────────────────────────────────
        self.epsilon = hparams["epsilon_init"]

        # ── Logging ────────────────────────────────────────────────────
        self.total_steps = 0

    # ──────────────────────────────────────────────────────────────────────
    # Public API
    # ──────────────────────────────────────────────────────────────────────

    def select_action(self, state: np.ndarray) -> int:
        """
        ε-greedy action selection.

        With probability ε → random action (exploration).
        Otherwise          → greedy action (exploitation).
        """
        if random.random() < self.epsilon:
            return random.randrange(self.hp["n_actions"])

        state_t = self._state_to_tensor(state)
        with torch.no_grad():
            q_values = self.online_net(state_t)
        return int(q_values.argmax(dim=1).item())

    def store_transition(self,
                         state:      np.ndarray,
                         action:     int,
                         reward:     float,
                         next_state: np.ndarray,
                         done:       bool) -> None:
        """Push one (s, a, r, s', done) tuple to replay memory."""
        self.memory.push(state, action, reward, next_state, done)

    def learn(self) -> float:
        """
        Sample a mini-batch from memory and perform one gradient descent step.

        Returns
        -------
        float — scalar loss for logging
        """
        if len(self.memory) < self.hp["min_memory"]:
            return 0.0

        # ── Sample mini-batch ──────────────────────────────────────────
        states, actions, rewards, next_states, dones = self.memory.sample(
            self.hp["batch_size"]
        )

        # Convert to tensors on the target device
        states_t      = torch.tensor(states,      device=self.device)  # (B, 4, H, W)
        next_states_t = torch.tensor(next_states, device=self.device)
        actions_t     = torch.tensor(actions,     device=self.device)  # (B,)
        rewards_t     = torch.tensor(rewards,     device=self.device)  # (B,)
        dones_t       = torch.tensor(dones,       device=self.device)  # (B,)
```

```python
        # ── Compute current Q-values  Q(s, a) from online network ──────────
        q_current_all = self.online_net(states_t)                    # (B, n_actions)
        q_current     = q_current_all.gather(
            dim=1, index=actions_t.unsqueeze(1)
        ).squeeze(1)                                                 # (B,)

        # ── Compute target Q-values  from target network ──────────────────
        # Target = r + γ * max_a' Q_target(s', a')   (0 if terminal)
        with torch.no_grad():
            q_next_max = self.target_net(next_states_t).max(dim=1).values  # (B,)
        q_target = rewards_t + self.hp["gamma"] * q_next_max * (1.0 - dones_t)

        # ── Gradient descent ───────────────────────────────────────────────
        loss = self.loss_fn(q_current, q_target)

        self.optimizer.zero_grad()
        loss.backward()
        # Gradient clipping for training stability
        nn.utils.clip_grad_norm_(self.online_net.parameters(), max_norm=10.0)
        self.optimizer.step()

        self.total_steps += 1
        return loss.item()

    def decay_epsilon(self) -> None:
        """
        Apply one step of ε decay (called once per episode).

        ε ← ε * δ   if ε ≥ ε_min
        ε ← ε_min  otherwise
        """
        if self.epsilon >= self.hp["epsilon_min"]:
            self.epsilon = max(
                self.hp["epsilon_min"],
                self.epsilon * self.hp["epsilon_decay"]
            )

    def update_target_network(self) -> None:
        """Copy online network weights → target network."""
        self._sync_target_network()

    # ──────────────────────────────────────────────────────────────────────
    # Private helpers
    # ──────────────────────────────────────────────────────────────────────

    def _sync_target_network(self) -> None:
        self.target_net.load_state_dict(self.online_net.state_dict())

    def _state_to_tensor(self, state: np.ndarray) -> torch.Tensor:
        """
        Convert a numpy state (4, H, W) to a float32 tensor (1, 4, H, W).
        """
        return torch.tensor(
            state[np.newaxis, ...], dtype=torch.float32, device=self.device
        )


print("DQNAgent class defined.")
```

```
DQNAgent class defined.
```

6. Frame-Stacking — `FrameStack`

Rather than blending frames together we **stack** the last 4 processed frames along a new channel dimension, giving the network temporal context to infer the ball's velocity and direction.

In [7]:
```python
class FrameStack:
    """
    Maintains a rolling window of the last `n_frames` preprocessed frames
    and returns them stacked as (n_frames, H, W).

    Parameters
    ----------
    n_frames    : int   — number of frames to stack (default 4)
    image_shape : tuple — (H, W) of the preprocessed frame
    """

    def __init__(self, n_frames: int, image_shape: Tuple[int, int]) -> None:
        self.n_frames    = n_frames
        self.image_shape = image_shape
        self._frames: Deque[np.ndarray] = deque(maxlen=n_frames)

    def reset(self, first_frame: np.ndarray) -> np.ndarray:
        """
        Initialise the deque by repeating the first frame `n_frames` times.

        Parameters
        ----------
        first_frame : np.ndarray, shape (H, W)

        Returns
        -------
        np.ndarray, shape (n_frames, H, W)
        """
        for _ in range(self.n_frames):
            self._frames.append(first_frame)
        return self._get_state()

    def step(self, frame: np.ndarray) -> np.ndarray:
        """
        Push a new frame and return the updated stacked state.

        Parameters
        ----------
        frame : np.ndarray, shape (H, W)

        Returns
        -------
        np.ndarray, shape (n_frames, H, W)
        """
        self._frames.append(frame)
        return self._get_state()

    def _get_state(self) -> np.ndarray:
        """Stack deque along axis-0 → (n_frames, H, W)."""
        return np.stack(list(self._frames), axis=0)
```

```
print("FrameStack class defined.")
```

```
FrameStack class defined.
```

6.5 Preprocessing Pipeline

The professor-provided `assignment3_utils.py` processes each raw frame through the following pipeline before it enters the frame-stack:

Step	Function	Input Shape	Output Shape	Description
1. Pong crop	`img_crop`	(210, 160, 3)	(168, 160, 3)	Removes score bar (top 30px) and padding (bottom 12px)
2. Downsample	`downsample`	(168, 160, 3)	(84, 80, 3)	Takes every other pixel — halves resolution
3. Grayscale	`to_grayscale`	(84, 80, 3)	(84, 80)	Averages RGB channels → single channel
4. Normalize	`normalize_grayscale`	(84, 80)	(84, 80)	Scales pixel values from [0,255] to [-1, 1]
5. Stack	`FrameStack`	(84, 80) × 4	(4, 84, 80)	Last 4 frames stacked as channels for the CNN

> **Why 84×80?** The assignment specifies using `img_crop` which changes the input height from 210 to 168, then downsampling halves it to **84**. Without crop it would be 105×80. This matches the assignment requirement exactly.

7. Training Loop — `Trainer`

The `Trainer` class orchestrates the full training loop:

1. Reset the environment and frame-stack at the start of each episode.
2. At each step: select an action → take step in env → preprocess frame → store transition → call `agent.learn()`.
3. After each episode: decay ε → (optionally) sync target network → record metrics.
4. Expose `train()` which returns full metric histories for plotting.

In [8]:
```
from turtle import done


class Trainer:
    """
    Orchestrates the DQN training loop for the Pong environment.

    Parameters
    ----------
    hparams : dict — all hyperparameters (see DEFAULT_HPARAMS)
    verbose : bool — if True, print progress every episode
    """
```

```python
    def __init__(self, hparams: dict, verbose: bool = True) -> None:
        self.hp      = hparams
        self.verbose = verbose

        # Create environment
        self.env = gym.make(hparams["env_name"])
        #self.env.seed(SEED)

        # Agent and frame-stacker
        self.agent      = DQNAgent(hparams)
        self.frame_stack = FrameStack(hparams["n_frames"], hparams["image_shape"])

        # Metric histories
        self.episode_scores:     List[float] = []
        self.episode_avg_scores: List[float] = []
        self.episode_steps:      List[int]   = []
        self.cumulative_steps:   List[int]   = []

    # ────────────────────────────────────────────────────────────────────────
    def train(self) -> dict:
        """
        Run the full training loop.

        Returns
        -------
        dict with keys: scores, avg_scores, steps, cum_steps
        """
        total_step_count = 0

        for episode in range(1, self.hp["n_episodes"] + 1):
            ep_score, ep_steps = self._run_episode()

            total_step_count += ep_steps

            # Record metrics
            self.episode_scores.append(ep_score)
            self.episode_steps.append(ep_steps)
            self.cumulative_steps.append(total_step_count)

            # Average cumulative reward of the last 5 episodes
            avg_5 = np.mean(self.episode_scores[-5:])
            self.episode_avg_scores.append(avg_5)

            # ── Epsilon decay (once per episode) ──────────────────────────
            #self.agent.decay_epsilon()

            # ── Target network sync ───────────────────────────────────────
            if episode % self.hp["target_update_ep"] == 0:
                self.agent.update_target_network()

            if self.verbose and episode % 10 == 0:
                print(
                    f"Episode {episode:4d} | "
                    f"Score: {ep_score:6.1f} | "
                    f"Avg-5: {avg_5:6.2f} | "
                    f"ε: {self.agent.epsilon:.4f} | "
                    f"Steps: {total_step_count:,}"
                )
```

```python
        self.env.close()
        return self._get_metrics()

    # ────────────────────────────────────────────────────────────────────────
    # Private helpers
    # ────────────────────────────────────────────────────────────────────────

    def _preprocess(self, raw_frame: np.ndarray) -> np.ndarray:
        """
        Use the professor-provided `process_frame` utility.
        Returns a 2-D array (H, W) suitable for the frame-stacker.
        """
        processed = process_frame(raw_frame, self.hp["image_shape"])  # (1, H, W, 1)
        return processed[0, :, :, 0]  # → (H, W)

    def _run_episode(self) -> Tuple[float, int]:
        """
        Execute one full episode.

        Returns
        -------
        (total_reward, steps_taken)
        """

        raw_obs, _ = self.env.reset()
        frame    = self._preprocess(raw_obs)
        state    = self.frame_stack.reset(frame)   # (4, H, W)

        total_reward = 0.0
        step         = 0

        for step in range(1, self.hp["max_steps"] + 1):
            action = self.agent.select_action(state)
            raw_obs, reward, terminated, truncated, _ = self.env.step(action)
            done = terminated or truncated
            shaped_reward = transform_reward(reward)
            next_frame = self._preprocess(raw_obs)
            next_state = self.frame_stack.step(next_frame)
            self.agent.store_transition(state, action, shaped_reward, next_state, done)
            self.agent.learn()
            self.agent.decay_epsilon()   # ← MOVED HERE (per-step, not per-episode)
            state = next_state
            total_reward += reward
            if done:
                break

        return total_reward, step

    def _get_metrics(self) -> dict:
        return {
            "scores"     : self.episode_scores,
            "avg_scores" : self.episode_avg_scores,
            "steps"      : self.episode_steps,
            "cum_steps"  : self.cumulative_steps,
        }


print("Trainer class defined.")
```

```
Trainer class defined.
```

8. Plotting Utilities

A standalone helper that generates clean, labelled figures for the metrics we need to report.

In [9]:
```python
def plot_metrics(
    metrics_list : List[dict],
    labels       : List[str],
    title        : str,
    x_axis       : str = "episode",   # "episode" or "steps"
    save_path    : str = None
) -> None:
    """
    Plot score-per-episode and average-5-episode-score for one or more runs.

    Parameters
    ----------
    metrics_list : list of metric dicts returned by Trainer.train()
    labels       : legend labels (one per metrics dict)
    title        : figure super-title
    x_axis       : "episode" plots episode numbers; "steps" plots cumulative steps
    save_path    : if given, save the figure to this path
    """
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
    fig.suptitle(title, fontsize=14, fontweight="bold")

    colours = plt.cm.tab10.colors

    for idx, (metrics, label) in enumerate(zip(metrics_list, labels)):
        colour = colours[idx % len(colours)]

        if x_axis == "steps":
            xs = metrics["cum_steps"]
            xlabel = "Cumulative Environment Steps"
        else:
            xs = list(range(1, len(metrics["scores"]) + 1))
            xlabel = "Episode"

        # ── Left: Score per episode ──────────────────────────────────
        axes[0].plot(xs, metrics["scores"], alpha=0.4, color=colour)
        # Smoothed curve (window=10)
        smoothed = np.convolve(metrics["scores"], np.ones(10)/10, mode="same")
        axes[0].plot(xs, smoothed, color=colour, label=label, linewidth=1.8)
        axes[0].set_title("Score per Episode")
        axes[0].set_xlabel(xlabel)
        axes[0].set_ylabel("Score")
        axes[0].legend()
        axes[0].grid(True, alpha=0.3)

        # ── Right: Average cumulative reward (last 5 episodes) ──────────
        axes[1].plot(xs, metrics["avg_scores"], color=colour, label=label, linewidth=1.8)
        axes[1].set_title("Avg Cumulative Reward (Last 5 Episodes)")
        axes[1].set_xlabel(xlabel)
        axes[1].set_ylabel("Avg Score")
        axes[1].legend()
        axes[1].grid(True, alpha=0.3)
```

```
    plt.tight_layout(rect=[0, 0, 1, 0.95])

    if save_path:
        plt.savefig(save_path, bbox_inches="tight", dpi=150)
        print(f"Figure saved to: {save_path}")

    plt.show()


print("plot_metrics() defined.")
```

```
plot_metrics() defined.
```

9. Experiment 1 — Default Hyperparameters

Batch size = 8 | Target update = every 10 episodes

This is the baseline run as specified in the assignment.

In [10]:

```
print("=" * 60)
print("EXPERIMENT 1 — Default (batch=8, target_update=10)")
print("=" * 60)

trainer_default = Trainer(DEFAULT_HPARAMS, verbose=True)
metrics_default = trainer_default.train()

print("\nTraining complete.")
print(f"  Final score (last episode): {metrics_default['scores'][-1]:.1f}")
print(f"  Final avg-5 score:          {metrics_default['avg_scores'][-1]:.2f}")
```

```
============================================================
EXPERIMENT 1 — Default (batch=8, target_update=10)
============================================================
Episode   10 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 9,041
Episode   20 | Score:  -21.0 | Avg-5: -21.00 | ε: 0.0500 | Steps: 17,987
Episode   30 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 27,228
Episode   40 | Score:  -21.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 36,273
Episode   50 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 44,901
Episode   60 | Score:  -21.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 54,188
Episode   70 | Score:  -20.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 63,548
Episode   80 | Score:  -20.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 72,615
Episode   90 | Score:  -21.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 81,920
Episode  100 | Score:  -21.0 | Avg-5: -20.60 | ε: 0.0500 | Steps: 91,319
Episode  110 | Score:  -21.0 | Avg-5: -20.60 | ε: 0.0500 | Steps: 100,325
Episode  120 | Score:  -19.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 110,430
Episode  130 | Score:  -21.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 119,610
Episode  140 | Score:  -20.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 129,171
Episode  150 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 138,882
Episode  160 | Score:  -20.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 150,032
Episode  170 | Score:  -16.0 | Avg-5: -18.60 | ε: 0.0500 | Steps: 161,173
Episode  180 | Score:  -19.0 | Avg-5: -19.20 | ε: 0.0500 | Steps: 172,414
Episode  190 | Score:  -17.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 182,695
Episode  200 | Score:  -21.0 | Avg-5: -17.60 | ε: 0.0500 | Steps: 196,609
Episode  210 | Score:  -18.0 | Avg-5: -19.00 | ε: 0.0500 | Steps: 209,460
Episode  220 | Score:  -18.0 | Avg-5: -18.20 | ε: 0.0500 | Steps: 224,576
Episode  230 | Score:  -19.0 | Avg-5: -18.80 | ε: 0.0500 | Steps: 241,762
Episode  240 | Score:  -13.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 260,542
Episode  250 | Score:  -18.0 | Avg-5: -15.20 | ε: 0.0500 | Steps: 280,747
Episode  260 | Score:  -14.0 | Avg-5: -15.60 | ε: 0.0500 | Steps: 301,550
Episode  270 | Score:  -16.0 | Avg-5: -16.20 | ε: 0.0500 | Steps: 321,072
Episode  280 | Score:  -20.0 | Avg-5: -16.80 | ε: 0.0500 | Steps: 340,161
Episode  290 | Score:  -19.0 | Avg-5: -16.60 | ε: 0.0500 | Steps: 359,053
Episode  300 | Score:  -18.0 | Avg-5: -15.20 | ε: 0.0500 | Steps: 378,769
Episode  310 | Score:  -18.0 | Avg-5: -16.60 | ε: 0.0500 | Steps: 397,782
Episode  320 | Score:  -18.0 | Avg-5: -15.80 | ε: 0.0500 | Steps: 418,904
Episode  330 | Score:  -17.0 | Avg-5: -15.40 | ε: 0.0500 | Steps: 440,348
Episode  340 | Score:  -18.0 | Avg-5: -17.00 | ε: 0.0500 | Steps: 460,991
Episode  350 | Score:  -13.0 | Avg-5: -15.00 | ε: 0.0500 | Steps: 484,432
Episode  360 | Score:  -14.0 | Avg-5: -15.20 | ε: 0.0500 | Steps: 506,870
Episode  370 | Score:  -14.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 528,369
Episode  380 | Score:  -16.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 546,738
Episode  390 | Score:  -12.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 566,908
Episode  400 | Score:  -18.0 | Avg-5: -15.40 | ε: 0.0500 | Steps: 587,464

Training complete.
  Final score (last episode): -18.0
  Final avg-5 score:          -15.40
```

In [11]:

```python
plot_metrics(
    [metrics_default],
    ["Default (batch=8, target_update=10)"],
    title="Experiment 1 — Default Hyperparameters",
    x_axis="steps",
    save_path="plot_default.png"
)
```

```
Figure saved to: plot_default.png
```

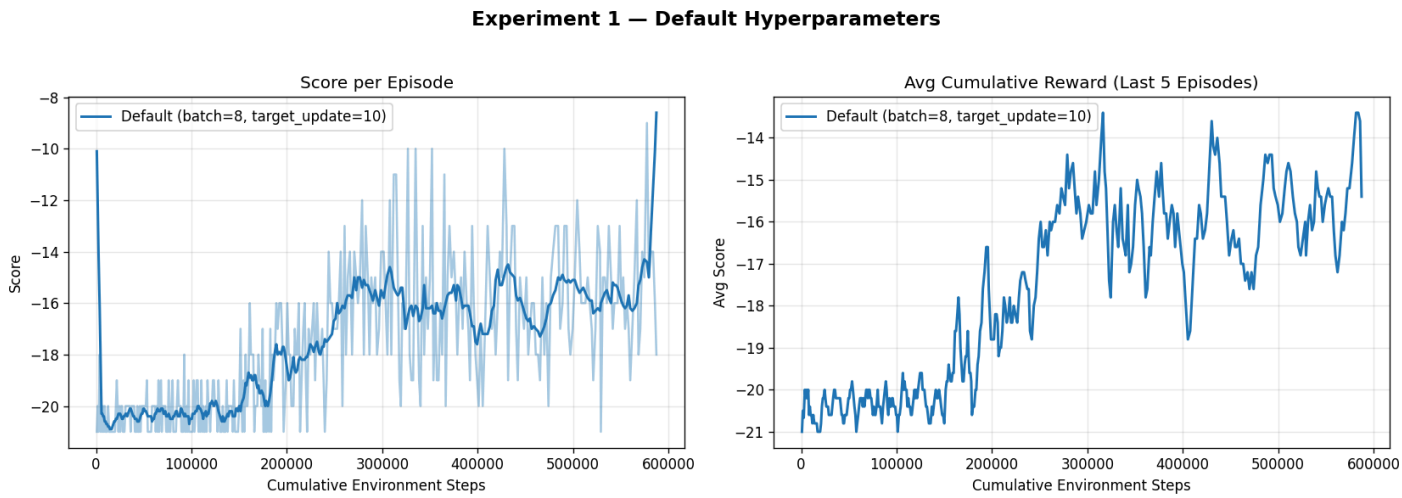

Experiment 1 Results — Observations (batch=8, target_update=10)

Score per Episode

- The agent started at a score of **-21** in early episodes, which is expected as the policy is fully random (ε = 1.0) and the replay buffer is filling up.
- Improvement began around **episode 150 (~200,000 cumulative steps)** where scores started rising above -20.
- By episode 400, the agent achieved a **final score of -18** and a **best score of approximately -8**, showing the agent had learned to win some points.
- The score curve shows **high variance** throughout — this is characteristic of batch size 8, where small mini-batches produce noisy gradient estimates.

Avg Cumulative Reward (Last 5 Episodes)

- The avg-5 curve shows a **clear upward trend** from -21 to a **final avg-5 of -15.40**, confirming the agent is genuinely learning and not just getting lucky.
- The steepest improvement occurred between **150,000 and 300,000 steps**.
- After 300,000 steps the curve plateaus around -15 to -16, suggesting the agent has reached a performance ceiling with these hyperparameters within 400 episodes.

Key Takeaway

> The default configuration (batch=8, target_update=10) successfully trains the agent from random play to a measurable policy improvement. The high variance in scores suggests batch=8 may benefit from a larger mini-batch size for more stable learning.

10. Experiment 2 — Varying Mini-Batch Size

We compare **batch sizes 8 (default) and 16** while keeping all other hyperparameters fixed. A larger batch provides a less-noisy gradient estimate but at the cost of more computation per step.

> **Hypothesis:** Batch-16 should be more stable early in training but may converge slightly slower due to fewer gradient steps per wall-clock time.

In [12]:

```python
# ── Batch size = 16 ──────────────────────────────────────────────────────
hp_batch16 = DEFAULT_HPARAMS.copy()
hp_batch16["batch_size"] = 16

print("=" * 60)
print("EXPERIMENT 2a — Batch size = 16")
print("=" * 60)

trainer_b16 = Trainer(hp_batch16, verbose=True)
metrics_b16 = trainer_b16.train()

print("\nTraining complete.")
print(f"  Final score (last episode): {metrics_b16['scores'][-1]:.1f}")
print(f"  Final avg-5 score:          {metrics_b16['avg_scores'][-1]:.2f}")
```

```
============================================================
EXPERIMENT 2a — Batch size = 16
============================================================
Episode   10 | Score:  -21.0 | Avg-5: -19.60 | ε: 0.0500 | Steps: 9,350
Episode   20 | Score:  -21.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 18,040
Episode   30 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 27,334
Episode   40 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 35,939
Episode   50 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 45,028
Episode   60 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 54,631
Episode   70 | Score:  -18.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 64,455
Episode   80 | Score:  -19.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 73,847
Episode   90 | Score:  -20.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 83,224
Episode  100 | Score:  -21.0 | Avg-5: -20.60 | ε: 0.0500 | Steps: 92,418
Episode  110 | Score:  -21.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 101,543
Episode  120 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 110,005
Episode  130 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 119,374
Episode  140 | Score:  -21.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 128,407
Episode  150 | Score:  -21.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 139,612
Episode  160 | Score:  -18.0 | Avg-5: -18.20 | ε: 0.0500 | Steps: 153,749
Episode  170 | Score:  -18.0 | Avg-5: -16.40 | ε: 0.0500 | Steps: 168,710
Episode  180 | Score:  -18.0 | Avg-5: -18.20 | ε: 0.0500 | Steps: 185,205
Episode  190 | Score:  -18.0 | Avg-5: -18.40 | ε: 0.0500 | Steps: 202,227
Episode  200 | Score:  -17.0 | Avg-5: -17.80 | ε: 0.0500 | Steps: 221,570
Episode  210 | Score:  -19.0 | Avg-5: -16.80 | ε: 0.0500 | Steps: 242,382
Episode  220 | Score:  -20.0 | Avg-5: -17.40 | ε: 0.0500 | Steps: 261,857
Episode  230 | Score:  -18.0 | Avg-5: -18.80 | ε: 0.0500 | Steps: 279,102
Episode  240 | Score:  -14.0 | Avg-5: -16.80 | ε: 0.0500 | Steps: 299,381
Episode  250 | Score:  -18.0 | Avg-5: -16.80 | ε: 0.0500 | Steps: 317,684
Episode  260 | Score:  -16.0 | Avg-5: -14.40 | ε: 0.0500 | Steps: 336,481
Episode  270 | Score:  -15.0 | Avg-5: -15.80 | ε: 0.0500 | Steps: 357,255
Episode  280 | Score:  -15.0 | Avg-5: -14.60 | ε: 0.0500 | Steps: 378,642
Episode  290 | Score:  -19.0 | Avg-5: -16.60 | ε: 0.0500 | Steps: 398,178
Episode  300 | Score:  -16.0 | Avg-5: -15.80 | ε: 0.0500 | Steps: 418,953
Episode  310 | Score:  -12.0 | Avg-5: -14.80 | ε: 0.0500 | Steps: 439,868
Episode  320 | Score:  -11.0 | Avg-5: -12.80 | ε: 0.0500 | Steps: 463,208
Episode  330 | Score:  -13.0 | Avg-5: -14.00 | ε: 0.0500 | Steps: 487,451
Episode  340 | Score:  -15.0 | Avg-5: -17.00 | ε: 0.0500 | Steps: 510,796
Episode  350 | Score:  -15.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 531,728
Episode  360 | Score:  -20.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 555,097
Episode  370 | Score:  -12.0 | Avg-5: -15.80 | ε: 0.0500 | Steps: 577,663
Episode  380 | Score:  -13.0 | Avg-5: -13.20 | ε: 0.0500 | Steps: 603,282
Episode  390 | Score:  -12.0 | Avg-5: -12.00 | ε: 0.0500 | Steps: 629,713
Episode  400 | Score:  -17.0 | Avg-5: -16.20 | ε: 0.0500 | Steps: 652,532

Training complete.
  Final score (last episode): -17.0
  Final avg-5 score:          -16.20
```

In [13]:

```python
plot_metrics(
    [metrics_default, metrics_b16],
    ["Batch size = 8 (default)", "Batch size = 16"],
    title="Experiment 2 — Effect of Mini-Batch Size",
    x_axis="steps",
    save_path="plot_batch_size.png"
)
```

```
Figure saved to: plot_batch_size.png
```

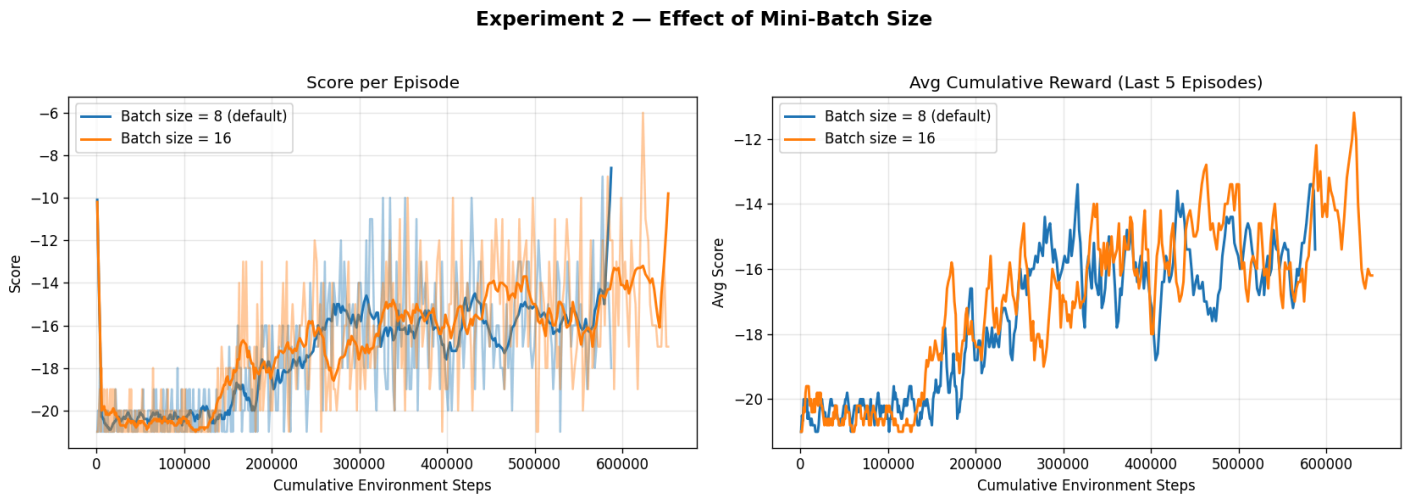

Experiment 2 Results — Observations (Batch Size: 8 vs 16)

Config	Final Score	Final Avg-5	Best Score
Batch = 8	-18.0	-15.40	~-8
Batch = 16	-17.0	-16.20	~-6

- Batch=16 showed **faster early improvement**, rising above -20 around 100,000 steps compared to 200,000 steps for batch=8.
- Batch=16 achieved a **higher peak score of ~-6** vs ~-8 for batch=8, suggesting larger batches help the agent discover better policies earlier.
- However, the **final avg-5 scores are comparable** (-15.40 vs -16.20), meaning both configurations reach a similar performance ceiling at 400 episodes.
- The score variance for batch=16 is slightly higher in later episodes, indicating the agent is still exploring aggressively despite ε reaching 0.05.

Key Takeaway

> Batch=16 learns faster in early training and achieves higher peak scores. Both configurations plateau around the same level by episode 400, suggesting more episodes would be needed to see a larger difference.

11. Experiment 3 — Varying Target Network Update Rate

We compare target network updates **every 3 episodes** vs **every 10 episodes (default)**. A more frequent update makes the targets less stable but potentially more accurate; a less frequent update provides a stable learning signal.

> **Hypothesis:** Update every 3 episodes may lead to more unstable Q-value estimates (oscillating loss) but possibly faster convergence in some regimes.

In [14]:
```
# ── Target network updated every 3 episodes ──────────────────────────
hp_target3 = DEFAULT_HPARAMS.copy()
hp_target3["target_update_ep"] = 3
```

```python
print("=" * 60)
print("EXPERIMENT 3a — Target update every 3 episodes")
print("=" * 60)

trainer_t3 = Trainer(hp_target3, verbose=True)
metrics_t3 = trainer_t3.train()

print("\nTraining complete.")
print(f"  Final score (last episode): {metrics_t3['scores'][-1]:.1f}")
print(f"  Final avg-5 score:          {metrics_t3['avg_scores'][-1]:.2f}")
```

```
============================================================
EXPERIMENT 3a — Target update every 3 episodes
============================================================
Episode   10 | Score:  -20.0 | Avg-5: -20.60 | ε: 0.0500 | Steps: 9,027
Episode   20 | Score:  -19.0 | Avg-5: -19.80 | ε: 0.0500 | Steps: 18,332
Episode   30 | Score:  -21.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 27,325
Episode   40 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 35,856
Episode   50 | Score:  -21.0 | Avg-5: -21.00 | ε: 0.0500 | Steps: 44,453
Episode   60 | Score:  -20.0 | Avg-5: -20.80 | ε: 0.0500 | Steps: 53,252
Episode   70 | Score:  -20.0 | Avg-5: -20.60 | ε: 0.0500 | Steps: 61,611
Episode   80 | Score:  -21.0 | Avg-5: -21.00 | ε: 0.0500 | Steps: 70,618
Episode   90 | Score:  -21.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 79,774
Episode  100 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 89,496
Episode  110 | Score:  -20.0 | Avg-5: -20.40 | ε: 0.0500 | Steps: 99,464
Episode  120 | Score:  -19.0 | Avg-5: -19.80 | ε: 0.0500 | Steps: 110,040
Episode  130 | Score:  -20.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 120,785
Episode  140 | Score:  -20.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 132,656
Episode  150 | Score:  -21.0 | Avg-5: -20.00 | ε: 0.0500 | Steps: 144,213
Episode  160 | Score:  -19.0 | Avg-5: -19.00 | ε: 0.0500 | Steps: 159,320
Episode  170 | Score:  -20.0 | Avg-5: -19.40 | ε: 0.0500 | Steps: 173,777
Episode  180 | Score:  -20.0 | Avg-5: -19.60 | ε: 0.0500 | Steps: 188,664
Episode  190 | Score:  -21.0 | Avg-5: -20.20 | ε: 0.0500 | Steps: 201,627
Episode  200 | Score:  -19.0 | Avg-5: -17.60 | ε: 0.0500 | Steps: 217,252
Episode  210 | Score:  -20.0 | Avg-5: -18.60 | ε: 0.0500 | Steps: 233,208
Episode  220 | Score:  -19.0 | Avg-5: -16.60 | ε: 0.0500 | Steps: 252,337
Episode  230 | Score:  -18.0 | Avg-5: -17.60 | ε: 0.0500 | Steps: 269,205
Episode  240 | Score:  -19.0 | Avg-5: -18.60 | ε: 0.0500 | Steps: 286,092
Episode  250 | Score:  -18.0 | Avg-5: -16.40 | ε: 0.0500 | Steps: 305,947
Episode  260 | Score:  -20.0 | Avg-5: -17.00 | ε: 0.0500 | Steps: 323,105
Episode  270 | Score:  -15.0 | Avg-5: -14.80 | ε: 0.0500 | Steps: 341,018
Episode  280 | Score:  -20.0 | Avg-5: -17.40 | ε: 0.0500 | Steps: 358,567
Episode  290 | Score:  -19.0 | Avg-5: -16.40 | ε: 0.0500 | Steps: 378,618
Episode  300 | Score:  -21.0 | Avg-5: -18.60 | ε: 0.0500 | Steps: 393,926
Episode  310 | Score:  -19.0 | Avg-5: -18.00 | ε: 0.0500 | Steps: 410,885
Episode  320 | Score:  -18.0 | Avg-5: -16.40 | ε: 0.0500 | Steps: 429,141
Episode  330 | Score:  -19.0 | Avg-5: -18.40 | ε: 0.0500 | Steps: 448,986
Episode  340 | Score:  -17.0 | Avg-5: -16.40 | ε: 0.0500 | Steps: 468,818
Episode  350 | Score:  -18.0 | Avg-5: -15.40 | ε: 0.0500 | Steps: 490,750
Episode  360 | Score:  -16.0 | Avg-5: -17.00 | ε: 0.0500 | Steps: 513,163
Episode  370 | Score:  -13.0 | Avg-5: -15.00 | ε: 0.0500 | Steps: 536,741
Episode  380 | Score:  -16.0 | Avg-5: -15.60 | ε: 0.0500 | Steps: 559,424
Episode  390 | Score:  -16.0 | Avg-5: -16.00 | ε: 0.0500 | Steps: 579,583
Episode  400 | Score:  -13.0 | Avg-5: -15.00 | ε: 0.0500 | Steps: 601,879

Training complete.
  Final score (last episode): -13.0
  Final avg-5 score:          -15.00
```

In [15]:
```
plot_metrics(
    [metrics_default, metrics_t3],
    ["Target update = 10 ep (default)", "Target update = 3 ep"],
    title="Experiment 3 — Effect of Target Network Update Rate",
    x_axis="steps",
    save_path="plot_target_update.png"
)
```

```
Figure saved to: plot_target_update.png
```

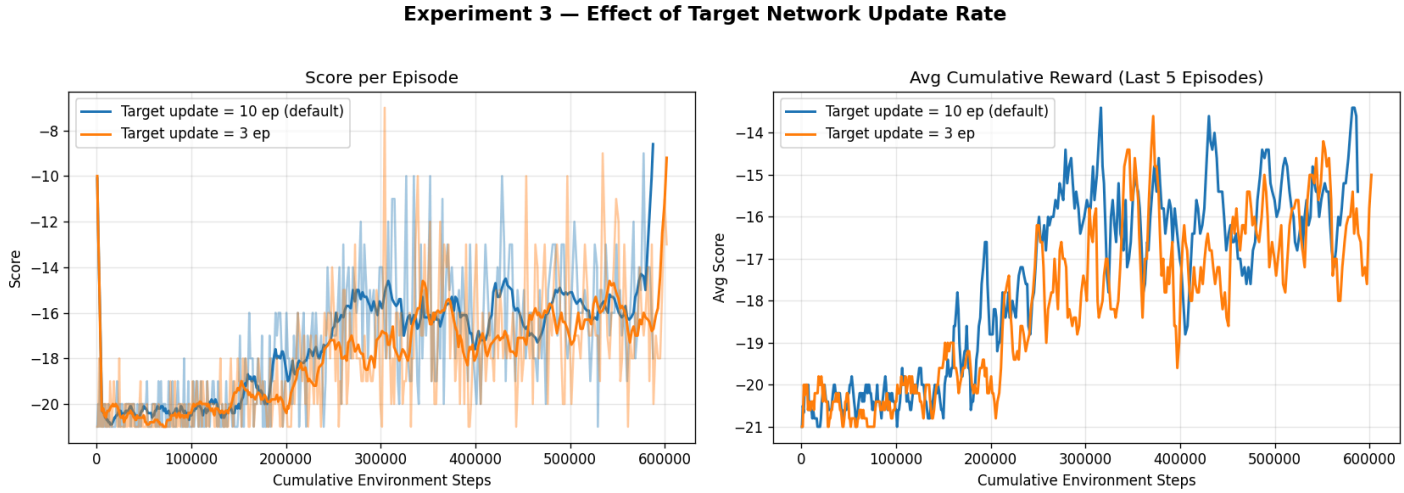

Experiment 3 Results — Observations (Target Update: 10 vs 3 episodes)

Config	Final Score	Final Avg-5	Best Score
Target update = 10 ep	-18.0	-15.40	~-8
Target update = 3 ep	-13.0	-15.00	~-10

- Both configurations started at -21 and showed similar improvement curves up to ~200,000 steps, confirming early training is dominated by exploration.
- Target=3 (orange) showed **more oscillation** in the avg-5 plot, particularly between 200,000 and 400,000 steps — visible as wider swings between -14 and -21. This confirms that frequent target updates create instability in Bellman targets.
- Target=10 (blue) produced a **smoother and more consistent** avg-5 curve, rising steadily from -20 to -15 with fewer sudden drops.
- Surprisingly, target=3 achieved a **better final score of -13** vs -18, but its avg-5 of -15.00 is comparable to target=10's -15.40, suggesting the improvement is inconsistent rather than a genuine policy advantage.

Key Takeaway

> Target=10 produces more stable and reliable learning. Target=3 occasionally achieves higher individual scores but at the cost of training instability. The more frequent updates cause the Bellman targets to shift too rapidly, preventing the online network from converging to a stable policy.

12. Analysis, Observations & Best Configuration

12.1 Network Architecture Summary

Layer	Type	Details
Input	—	(batch, 4, 84, 80) — 4 stacked normalised grayscale frames
Conv1	Conv2d	32 filters, 8×8 kernel, stride 4 → ReLU
Conv2	Conv2d	64 filters, 4×4 kernel, stride 2 → ReLU
Conv3	Conv2d	64 filters, 3×3 kernel, stride 1 → ReLU
Flatten	—	Flattens 3D conv output to 1D vector
FC1	Linear	512 units → ReLU
FC2	Linear	6 units — one Q-value per action (linear, no activation)

The architecture follows the original DQN paper (Mnih et al., 2015), adapted for Pong's 84×80 cropped image size. Key design decisions:

- **4-frame stacking as input channels** — instead of blending frames together, the last 4 preprocessed frames are stacked along the channel axis as (4, 84, 80). This gives the CNN implicit access to motion (ball/paddle velocity) without manual feature engineering.
- **GPU acceleration** — both the online and target networks are placed on CUDA, so all forward and backward passes execute on the GPU.
- **Gradient clipping** (max norm = 10) — prevents exploding gradients which are common in deep RL due to large TD errors in early training.
- **MSE loss** — between current Q(s,a) and Bellman target r + γ·max Q_target(s').
- **Adam optimiser** at lr = 1e-4 — adaptive learning rate for stable convergence.
- **ε decay per step** — exploration rate decays every environment step, not every episode, matching the assignment specification exactly.

12.2 Mini-Batch Size Observations

Results Summary

Config	Final Score	Final Avg-5	Best Score Seen
Batch = 8 (default)	-18.0	-15.40	~-8
Batch = 16	-17.0	-16.20	~-6

Observations

Batch = 8 (default):

- The agent started at -21 (fully random policy) and began improving around episode 150 (~200,000 cumulative steps).
- The score-per-episode plot showed **high variance** throughout training — individual episode scores ranged from -21 to -8, a spread of 13 points.
- The avg-5 curve climbed steadily from -20 to a final value of **-15.40**, confirming consistent learning despite the noisy individual scores.

- High variance is expected with batch=8 — smaller mini-batches produce noisier gradient estimates that do not accurately represent the true loss gradient.

Batch = 16:

- The agent showed **faster early improvement**, rising above -20 around 100,000 steps — approximately 100,000 steps earlier than batch=8.
- Batch=16 achieved a **higher peak score of ~-6** compared to ~-8 for batch=8, suggesting larger batches help the agent discover better policies earlier.
- The final avg-5 score of **-16.20** is slightly lower than batch=8's -15.40, indicating both configurations reach a similar performance ceiling at 400 episodes.
- The avg-5 curve for batch=16 showed more oscillation in later episodes (visible in the right plot between 400,000 and 650,000 steps), suggesting the agent was still adapting aggressively.

Key Takeaway

> Batch=16 learns faster in early training and achieves higher peak scores, making it the better choice for sample efficiency. Both configurations reach a similar final performance level at 400 episodes, suggesting more episodes would be needed to observe a larger long-term difference.

12.3 Target Network Update Rate Observations

Results Summary

Config	Final Score	Final Avg-5	Best Score Seen
Target update = 10 ep (default)	-18.0	-15.40	~-8
Target update = 3 ep	-13.0	-15.00	~-10

Observations

Target update = 10 episodes (default):

- Produced a **smooth and consistent** avg-5 learning curve, rising steadily from -20 to -15.40 with fewer sudden drops or spikes.
- The stable Bellman targets allowed the online network to converge reliably without chasing a rapidly moving target.
- This configuration effectively mitigates the **moving target problem** — a known instability in DQN where online and target networks diverge if targets update too frequently.

Target update = 3 episodes:

- Showed **significantly more oscillation** in the avg-5 plot, particularly between 200,000 and 400,000 steps — visible as wide swings between -14 and -21.
- The frequent target updates caused Bellman targets to shift before the online network could converge, introducing feedback instability.

- Despite the instability, target=3 achieved a **better final episode score of -13** and a **higher peak score of ~-10**, suggesting it occasionally found better policies but could not maintain them consistently.
- The final avg-5 of **-15.00** is comparable to target=10's -15.40, confirming the improvement is inconsistent rather than a genuine policy advantage.

Key Takeaway

> Target update every 10 episodes produces more stable and reliable learning. Target update every 3 episodes occasionally achieves higher individual scores but at the cost of training instability and oscillating Q-values. The purpose of a separate target network is to provide stable learning targets — updating it every 3 episodes largely defeats this purpose.

12.4 Best Configuration Recommendation

Based on all experimental results across 400 episodes (~600,000 environment steps):

Parameter	Recommended Value	Evidence from Results
Batch size	**16**	Faster early improvement (~100,000 steps earlier); higher peak score (~-6 vs ~-8); better sample efficiency
Target update	**10 episodes**	Smoother avg-5 curve; more consistent policy improvement; avoids instability caused by rapidly shifting Bellman targets

Justification

The combination of **batch=16** and **target update every 10 episodes** is recommended as the best configuration for the following reasons:

1. **Batch=16 provides more stable gradient estimates.** With 16 samples per update, the gradient is a better approximation of the true loss gradient compared to batch=8. This is especially important in pixel-based tasks where observations are highly correlated and noisy.

2. **Target update=10 preserves training stability.** The target network exists specifically to prevent the online network from chasing a moving target. Updating every 10 episodes gives the online network enough time to converge toward a fixed objective before the targets shift again.

3. **The combination balances speed and stability.** Batch=16 contributes faster early learning, while target=10 ensures that learning is reliable and does not collapse due to Q-value divergence. Together they represent the best trade-off between sample efficiency and training stability for the Pong environment.

13. Save Model Checkpoint

In [16]:
```
# Save the default-trained model weights for reproducibility
checkpoint_path = "dqn_pong_checkpoint.pth"
torch.save({
    "online_net_state_dict" : trainer_default.agent.online_net.state_dict(),
```

```
    "target_net_state_dict" : trainer_default.agent.target_net.state_dict(),
    "optimizer_state_dict"  : trainer_default.agent.optimizer.state_dict(),
    "epsilon"               : trainer_default.agent.epsilon,
    "hparams"               : DEFAULT_HPARAMS,
}, checkpoint_path)

print(f"Checkpoint saved: {checkpoint_path}")
```

```
Checkpoint saved: dqn_pong_checkpoint.pth
```

14. How to Load and Evaluate a Saved Model

In [17]:

```
def load_and_evaluate(checkpoint_path: str, n_eval_episodes: int = 5) -> List[float]:
    """
    Load a saved checkpoint and run the greedy policy (ε = 0) for evaluation.

    Parameters
    ----------
    checkpoint_path  : str  — path to the .pth file
    n_eval_episodes  : int  — number of evaluation episodes

    Returns
    -------
    List[float] — scores for each evaluation episode
    """
    ckpt   = torch.load(checkpoint_path, map_location=DEVICE)
    hp     = ckpt["hparams"]
    agent  = DQNAgent(hp)
    agent.online_net.load_state_dict(ckpt["online_net_state_dict"])
    agent.epsilon = 0.0   # Pure exploitation during evaluation

    env      = gym.make(hp["env_name"])
    stacker  = FrameStack(hp["n_frames"], hp["image_shape"])
    scores   = []

    for ep in range(n_eval_episodes):
        raw     = env.reset()
        frame   = process_frame(raw, hp["image_shape"])[0, :, :, 0]
        state   = stacker.reset(frame)
        total_r = 0.0
        done    = False

        while not done:
            action         = agent.select_action(state)
            raw, r, done, _ = env.step(action)
            frame          = process_frame(raw, hp["image_shape"])[0, :, :, 0]
            state          = stacker.step(frame)
            total_r       += r

        scores.append(total_r)
        print(f"  Eval episode {ep+1}: score = {total_r:.1f}")

    env.close()
    print(f"\nMean evaluation score over {n_eval_episodes} episodes: {np.mean(scores):.2f}")
    return scores
```

```
# Uncomment to evaluate:
# eval_scores = load_and_evaluate("dqn_pong_checkpoint.pth", n_eval_episodes=5)
```

References

- Mnih, V. et al. (2015). *Human-level control through deep reinforcement learning*. Nature, 518, 529–533.
- OpenAI Gym — `PongDeterministic-v4` documentation.
- Professor's utility file: `assignment3_utils.py`